

Project 1: Threads

Preliminaries

Fill in your name and email address.

Tong WU 2200013212@stu.pku.edu.cn

If you have any preliminary comments on your submission, notes for the TAs, please give them here.

My local test result:

```
pass tests/threads/alarm-single
pass tests/threads/alarm-multiple
pass tests/threads/alarm-simultaneous
pass tests/threads/alarm-priority
pass tests/threads/alarm-zero
pass tests/threads/alarm-negative
pass tests/threads/priority-change
pass tests/threads/priority-donate-one
pass tests/threads/priority-donate-multiple
pass tests/threads/priority-donate-multiple2
pass tests/threads/priority-donate-nest
pass tests/threads/priority-donate-sema
pass tests/threads/priority-donate-lower
pass tests/threads/priority-fifo
pass tests/threads/priority-preempt
pass tests/threads/priority-sema
pass tests/threads/priority-condvar
pass tests/threads/priority-donate-chain
pass tests/threads/mlfqs-load-1
pass tests/threads/mlfqs-load-60
pass tests/threads/mlfqs-load-avg
pass tests/threads/mlfqs-recent-1
pass tests/threads/mlfqs-fair-2
pass tests/threads/mlfqs-fair-20
pass tests/threads/mlfqs-nice-2
pass tests/threads/mlfqs-nice-10
pass tests/threads/mlfqs-block
All 27 tests passed.
```

Please cite any offline or online sources you consulted while preparing your submission, other than the Pintos documentation, course text, lecture notes, and course staff.

Alarm Clock

DATA STRUCTURES

A1: Copy here the declaration of each new or changed struct or struct member, global or static variable, typedef, or enumeration. Identify the purpose of each in 25 words or less.

```
1 struct thread_sleep {
2     // Keep track of threads' time to sleep
3     struct thread *t; // The ptr to the sleeping thread
4     int64_t wakeup_time; // The time(tick) to wake up the thread
5     struct list_elem elem; // To use it in a Pintos list
6 };
7
8 static struct list thread_sleep_list; // List of thread_sleep
```

ALGORITHMS

A2: Briefly describe what happens in a call to `timer_sleep()`, including the effects of the timer interrupt handler.

In my design, first we should check if the ticks to sleep is not positive and disable interruption to avoid race conditions. Then we add the thread calling `timer_sleep()` to the global list of sleeping threads, recording the time to wake it up. Here we use `list_insert_ordered()` to ensure the list is in order with respect to the wake up time. Finally, we call `thread_block()` to keep the thread away from the ready queue.

When the timer interrupts, the timer interrupt handler is called. It will first increase the system tick by 1 then check the sleeping list from front to end whether it is time to wake up sleeping threads(i.e. removing them from sleeping list and adding to the ready queue). Also, to enable preemption, `intr_yield_on_return()` is called if at least one thread is waken up.

A3: What steps are taken to minimize the amount of time spent in the timer interrupt handler?

we use `list_insert_ordered()` to ensure the list is in order with respect to the wake up time. As a result, in `timer_interrupt()` we no longer need to sort the list or check every thread's wake up time. Instead, whenever we encounter the first thread that still needs sleeping, we can directly 'break' from the checking loop.

SYNCHRONIZATION

A4: How are race conditions avoided when multiple threads call `timer_sleep()` simultaneously?

When `timer_sleep()` is called, interrupt is disabled before being added to the list of sleeping threads.

A5: How are race conditions avoided when a timer interrupt occurs during a call to `timer_sleep()`?

When `timer_sleep()` is called, interrupt is disabled before being added to the list of sleeping threads.

RATIONALE

A6: Why did you choose this design? In what ways is it superior to another design you considered?

At first I considered simply adding a new thread to the sleeping list when it calls `sleep()`, and check all threads' wake up time when handling interruption of the timer. However the current implementation is more efficient, considering that we can only check a few threads in every interruption.

Also, not every thread will call `sleep()`, so I maintain the list of sleeping threads in `timer.c`, rather than adding new member variables to `struct thread`.

Priority Scheduling

DATA STRUCTURES

B1: Copy here the declaration of each new or changed struct or struct member, global or static variable, typedef, or enumeration. Identify the purpose of each in 25 words or less.

I add `priority` and `elem` in `struct lock`.

```
1  /** Lock. */
2  struct lock
3  {
4      ... // Previous variables
5      int priority;           /**< Priority of the lock, equal to the
6      struct list_elem elem;  /**< For lock list. */
7  };
```

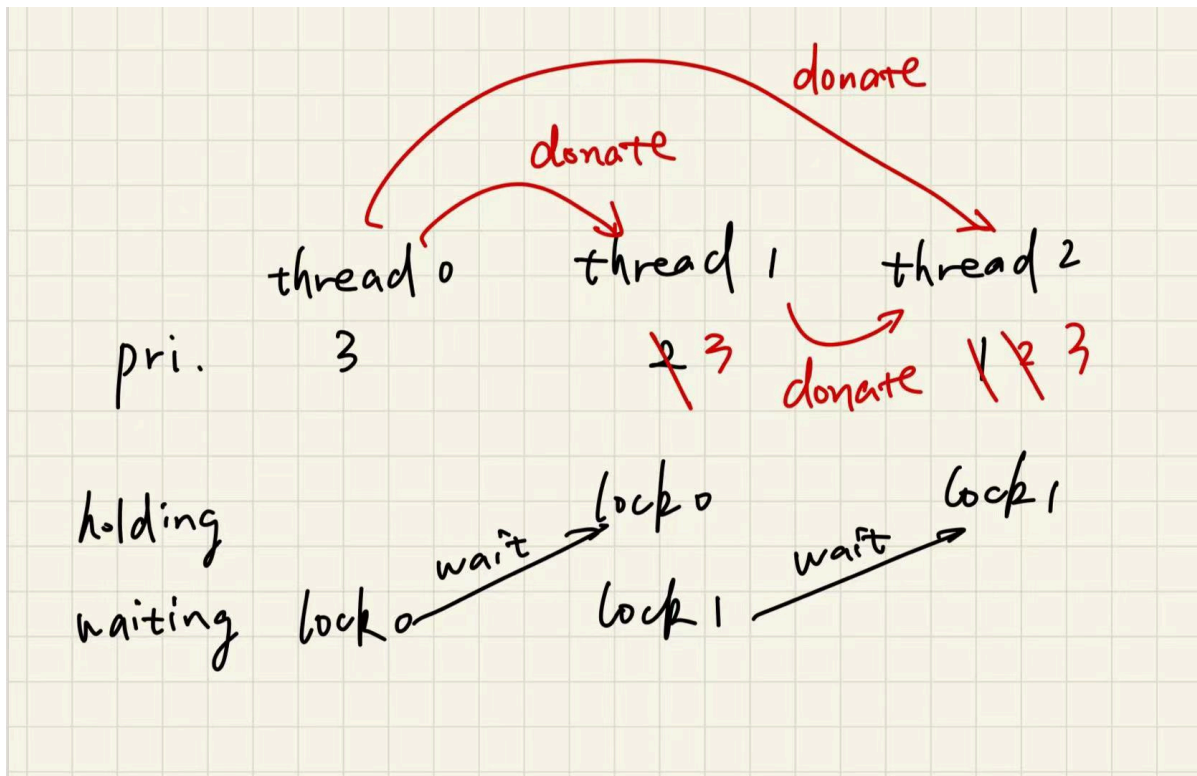
Also, I add these struct members to `struct thread`:

```
1  /* Utility for priority donation */
2      int old_priority;           /**< Original priority before donation
3  */
4      struct list holding_locks;  /**< Locks held by the thread. */
5      struct lock *waiting_lock;  /**< Lock that the thread is waiting
6  for. */
7      // Every thread can only acquire 1 lock, but holding >1 locks.
```

B2: Explain the data structure used to track priority donation.
Use ASCII art to diagram a nested donation. (Alternately, submit a .png file.)

Every thread can only acquire 1 lock, but can hold >1 locks. As a result, I use a `waiting_lock` member of each thread to keep track of the lock they're waiting for, and a list `holding_locks` to keep track of the lock they're holding. Also, `old_priority` is recorder, so that when a thread no more receives priority donation, its priority can be set back as before the donation.

A potential nested donation:



ALGORITHMS

B3: How do you ensure that the highest priority thread waiting for a lock, semaphore, or condition variable wakes up first?

- Semaphore: When `sema_up()` is called, I use `list_max()` to acquire the waiting thread with highest priority and unblock it first.
- Lock: In Pintos, lock is a specialization of a semaphore with an value ≤ 1 and only one holder, so all we need to do is ensure `sema_up()` is correctly implemented.
- Cond. var: The idea is similar to semaphore. When `condvar_signal()` is called, I compare different semaphores based on their waiter's highest priority and wake up the most prioritized waiter first.

B4: Describe the sequence of events when a call to `lock_acquire()` causes a priority donation. How is nested donation handled?

Suppose thread 1, 2, 3 has priority 1, 2, 3. Thread 3 calls `lock_acquire()` but thread 2 has higher priority than thread 1. So thread 3 donates its priority to thread 1 and now thread 1 has a priority of 3, which is higher than thread 2, so that it can successfully release the lock and let thread 3 run.

Nested donation is handled in `lock_acquire()`. When a thread is trying to acquire a lock, it will potentially recursively donate priority. It may donate to the holder of the lock it is waiting for, and to the holder of another lock that the holder of the lock it is waiting for is waiting for, and so on.

B5: Describe the sequence of events when `lock_release()` is called on a lock that a higher-priority thread is waiting for.

`lock_release()` first resets the relevant information, such as the lock's holder and the thread's `holding_locks`. Then it resets the current thread's priority to that before the lock is acquired by the current thread. Finally, `thread_yield()` is called to schedule the waiting thread with higher priority.

SYNCHRONIZATION

B6: Describe a potential race in `thread_set_priority()` and explain how your implementation avoids it. Can you use a lock to avoid this race?

Suppose a thread is lowering its priority, suddenly it's interrupted by the timer, and another thread is scheduled to run, which donates priority to the thread. However, the previous thread has decided to lower its priority, making the priority donation worthless.

I think that the code that I submitted did not avoid the potential race. I may disable the interruption during the function to avoid this race.

RATIONALE

B7: Why did you choose this design? In what ways is it superior to another design you considered?

Actually I only figure out the current design :(

I think one benefit of the current design is that the holder member of a lock and `holding_locks` and `waiting_lock` members of a thread enable us to conveniently recursively find the relationship between threads and locks, so that we can modify their information and priority without trouble.

Advanced Scheduler

DATA STRUCTURES

C1: Copy here the declaration of each new or changed struct or struct member, global or static variable, typedef, or enumeration. Identify the purpose of each in 25 words or less.

In `fixed-point.h`

```
1 typedef int fixed_point;
2 static int f = 16384; // 2^14
```

In `thread.c`

```
1 /** Global load_avg for mlfqs */
2 static fixed_point load_avg;
```

Adding new members to thread struct:

```
1 /* Values for mlfqs */
2 int nice;                               /**< Nice value for mlfqs. */
3 fixed_point recent_cpu;                 /**< Recent CPU for mlfqs. */
```

ALGORITHMS

C2: How is the way you divided the cost of scheduling between code inside and outside interrupt context likely to affect performance?

1. Updating `recent_cpu` every tick in the timer handler is executed only once and takes little time.

2. Updating `recent_cpu` and `load_avg` takes little time, too, because they are updated in this way once in a second.
3. Updating `priority` every 4 ticks is time-consuming, as we have to go through every thread.

RATIONALE

C3: Briefly critique your design, pointing out advantages and disadvantages in your design choices. If you were to have extra time to work on this part of the project, how might you choose to refine or improve your design?

- Advantages:
 - A simply implementation without multi-level queues is easy and straightforward.
- Disadvantages:
 - The implementation is not efficient enough, because I have to go through all threads whenever I decide the next thread to run.

I may improve by defining 64 queues with different priority level so that I can efficiently find the next thread with highest priority to run.

C4: The assignment explains arithmetic for fixed-point math in detail, but it leaves it open to you to implement it. Why did you decide to implement it the way you did? If you created an abstraction layer for fixed-point math, that is, an abstract data type and/or a set of functions or macros to manipulate fixed-point numbers, why did you do so? If not, why not?

I use `typedef` to alias `int32` as `fixed_point` so that it is efficient and explicit enough.

If I create an abstraction layer, it may be more explicit and provide better interface, whiling sacrificing some efficiency.